

Direct Preference Optimization was proposed by [1]. A follow-up idea, reward-free contrastive preference optimization, is described in [?].

References

- [1] Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D. Manning, Stefano Ermon, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. *arXiv preprint arXiv:2305.18290*, 2023.